

# Task 1.1 Define the problem

What exactly are we trying to solve with sentiment recognition?

1. In sentiment recognition, we are trying to identify and understand the emotional tone expressed in text, speech, images or any other data format (images in the assignment's case). The problem it solves is determining whether the expressed sentiment is positive, negative, neutral or more detailed such as anger, disgust, fear, happy, pain or sad, as given in the dataset we are working on in the assignment.
2. The main goal is to nullify the gap between human emotions and machine understanding, so the softwares can -
  - Interpret the sentiment of the data expressed.
  - Classify the feedbacks, reviews or conversations automatically
  - Provide insights for decision making

What are the desired inputs and outputs of the system?

1. The desired inputs are -
  - Raw dataset that contains the six folders containing the images showing that particular emotion.
  - So, the system will take as input an individual image from one of these folders.
2. The desired outputs are -
  - A predicted label that classifies the input into one of the six emotions.

Clearly describe the target classes (Happy, Sad, Fear, Pain, Anger, Disgust)

The target class in the assignment represent six distinct human emotions, each corresponding to a folder in the dataset. These are the categories that the CNN model will learn to recognize and classify from the input data:

1. Happy - Represents positive emotional state, generally expressed by a smile, joyful tone or cheerful expressions.
2. Sad - Represents negative emotional state, shown through crying or low energy.
3. Fear - Reflects a state of anxiety or being threatened, visible in widened eyes, tense body language or fearful tone.
4. Pain - Represents discomfort or suffering, often shown through facial expression or strained expressions.
5. Anger - Expresses frustration or aggression, commonly visible in frowns or glaring eyes.
6. Disgust - Represents aversion or dislike, often seen in wrinkled noses, curled lips or rejecting gestures.

For the CNN, these six classes form the output layer, where the model predicts a probability distribution across all classes and assigns the input to the class with the

highest probability.

Why might misclassification be problematic?

Misclassification in sentiment recognition can be problematic because:

- Loss of meaning - EMotions carry critical information about the customers or humans in general. Predicting anger as happy completely changes the intended meaning.
- Poor decision-making - Many applications (e.g. healthcare, customer service, security) rely on emotion detection. Misclassification could lead to wrong actions or responses.
- User dissatisfaction - If a system malfunctions, users may lose trust in the organisation or the vision of the organisation.
- Context sensitivity - Some emotions, like pain or happy, are two completely different emotions and misclassifying them could have serious consequences (e.g. in medical or safety contexts).

```
In [1]: import random
import numpy as np
import torch
import os

def set_seed(seed=42):
    random.seed(seed)                      # Python random
    np.random.seed(seed)                    # NumPy
    torch.manual_seed(seed)                # PyTorch CPU
    torch.cuda.manual_seed(seed)           # PyTorch GPU (single-GPU)
    torch.cuda.manual_seed_all(seed)       # PyTorch GPU (multi-GPU)

    # For deterministic behaviour (may slow training a bit)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

    os.environ["PYTHONHASHSEED"] = str(seed)

set_seed(42)
```

## Task 1.2 Make a plan

**Note - I have downloaded the dataset locally also, but as the dataset is very huge, it would take time to upload the dataset. So, let's use kaggle API.**

```
In [2]: !pip -q install kaggle
```

```
In [3]: from google.colab import files
files.upload()
```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

```
Out[3]: {'kaggle.json': b'{"username":"jp1611","key":"a25b9c0828d0bb54474f3659f71b79ec"}'}
```

```
In [4]: !mkdir -p ~/.kaggle  
!cp kaggle.json ~/.kaggle/  
!chmod 600 ~/.kaggle/kaggle.json
```

```
In [5]: !kaggle datasets download -d yousefmohamed20/sentiment-images-classifier -p /  
  
Dataset URL: https://www.kaggle.com/datasets/yousefmohamed20/sentiment-images-  
-classifier  
License(s): apache-2.0  
Downloading sentiment-images-classifier.zip to /content/data  
89% 101M/114M [00:00<00:00, 235MB/s]  
100% 114M/114M [00:00<00:00, 252MB/s]
```

```
In [6]: # Loading the dataset  
import os  
from torchvision import transforms  
from torchvision.datasets import ImageFolder  
from torch.utils.data import DataLoader  
  
data_root = "/content/data/6 Emotions for image classification"  
  
transform = transforms.Compose([  
    transforms.Resize((224, 224)),  
    transforms.ToTensor(),  
    transforms.Normalize(mean = [0.485, 0.456, 0.406], std = [0.229, 0.224,  
])  
  
dataset = ImageFolder(root = data_root, transform = transform)
```

```
In [7]: import matplotlib.pyplot as plt  
import numpy as np  
import torchvision  
  
data_loader = DataLoader(dataset, batch_size = 16, shuffle = True)  
  
images, labels = next(iter(data_loader))  
  
grid = torchvision.utils.make_grid(images, nrow = 4, padding = 2)  
  
# Convert from tensor (C, H, W) → (H, W, C) for plotting  
npimg = grid.numpy().transpose((1, 2, 0))  
  
# Undo normalization (so colors look right)  
mean = np.array([0.485, 0.456, 0.406])  
std = np.array([0.229, 0.224, 0.225])  
npimg = std * npimg + mean # reverse of the normalization  
npimg = np.clip(npimg, 0, 1)  
  
plt.figure(figsize = (8, 8))  
plt.imshow(npimg)  
plt.axis('off')  
plt.show()
```

```
print ("Labels: ", [dataset.classes[label] for label in labels])
```



Labels: ['happy', 'sad', 'fear', 'sad', 'happy', 'disgust', 'happy', 'disgust', 'sad', 'disgust', 'sad', 'disgust', 'sad', 'disgust', 'fear', 'anger']

```
In [8]: # Total number of images and images per label
from collections import Counter
print(f"Total images in the dataset are {len(dataset)}")

class_counts = Counter(dataset.targets)
for idx, count in class_counts.items():
    print(f"Class {dataset.classes[idx]}: {count} images")
```

Total images in the dataset are 1198  
Class anger: 214 images  
Class disgust: 199 images  
Class fear: 163 images  
Class happy: 230 images  
Class pain: 168 images  
Class sad: 224 images

1. Given the total number of images in the dataset, how would you determine an appropriate train/validation/test split ratio? Justify your choice?  
Traditionally, there are three splits, 80/10/10, 70/15/15 or 60/20/20 and the choice of the split depends on the size of the dataset.

As our dataset has 1198, it is a small dataset and if we use 80/10/10, our test and validation sets will have only 120 images, which are too small to properly evaluate.

If we use 70/15/15 split, our test and validation sets will have 180 images which is much better compared to 80/10/10.

If we select 60/20/20, our train set will have only 720 images, which are not sufficient for good training of the model.

Hence, the optimal choice is to use the 70/15/15 split.

2. Using quantitative analysis (e.g., histograms, percentages), evaluate whether the dataset is balanced across the six emotion classes. What evidence indicates underrepresentation, and how might this influence learning dynamics?

From the counts of each class -

3. Class anger: 214 images
4. Class disgust: 199 images
5. Class fear: 163 images
6. Class happy: 230 images
7. Class pain: 168 images
8. Class sad: 224 images

Hence,

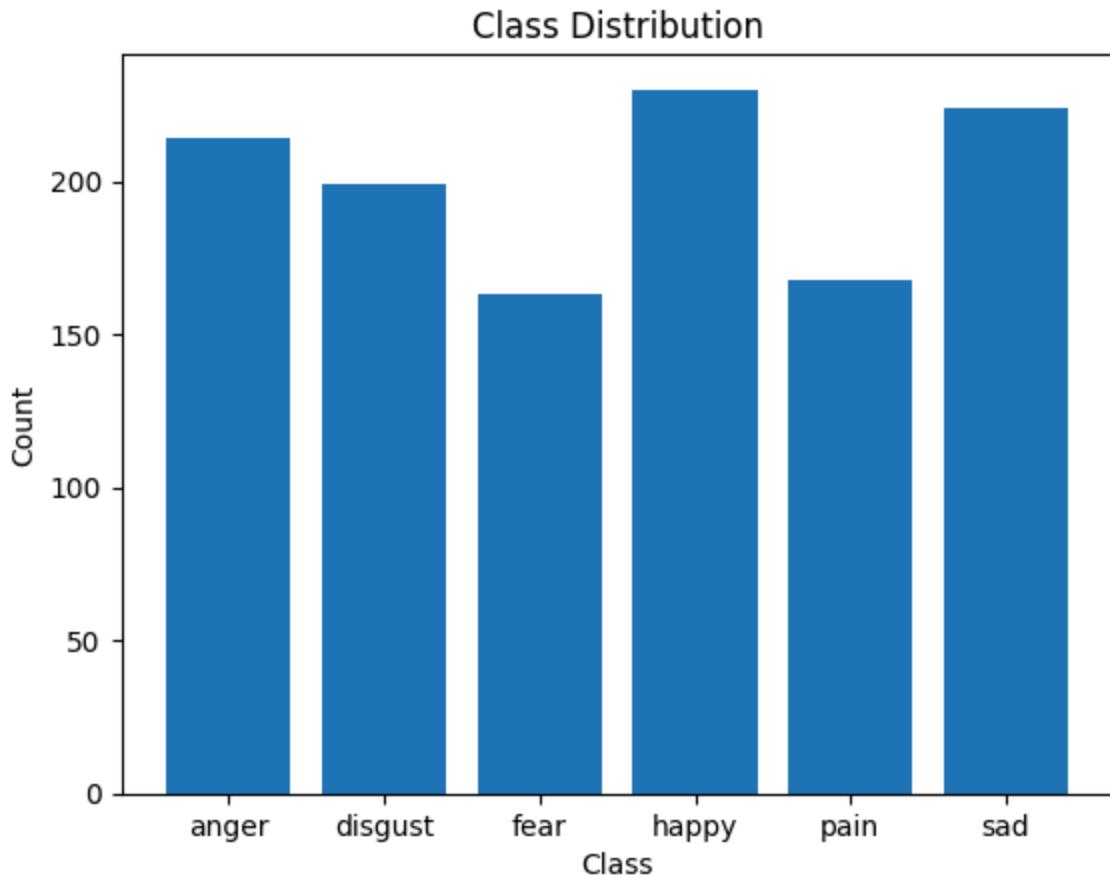
- anger:  $214 / 1198 \approx 17.9\%$
- disgust:  $199 / 1198 \approx 16.6\%$
- fear:  $163 / 1198 \approx 13.6\%$
- happy:  $230 / 1198 \approx 19.2\%$
- pain:  $168 / 1198 \approx 14.0\%$
- sad:  $224 / 1198 \approx 18.7\%$

The labels **fear** (13.6%) and **bold** (19.2%) are underrepresented.

**happy** label is overrepresented.

```
In [9]: labels = ["anger", "disgust", "fear", "happy", "pain", "sad"]
class_counts = [214, 199, 163, 230, 168, 224]

plt.bar(labels, class_counts)
plt.xlabel("Class")
plt.ylabel("Count")
plt.title("Class Distribution")
plt.show()
```



For the underrepresented classes, the model may see fewer examples and hence, it might fail to learn their features well leading to higher misclassification rates for these classes.

For overrepresented classes, the CNN may produce bias predictions towards these classes, since they appear more often.

To overcome the imbalance, we can use techniques such as **data augmentation** (oversampling, flipping, rotations, etc.) for minority classes.

3. From a theoretical standpoint, explain how class imbalance can bias gradient updates during training. Which corrective techniques (e.g., reweighting loss functions, oversampling, data augmentation) could be applied, and why?

During the training, the majority classes dominate the gradient updates because they contribute more samples to the loss.

This means -

- The CNN weights are optimized more for **majority class**.
- Minority classes contribute less frequently, so their error signal is weaker and often gets “drowned out.”

As a result, the model biases predictions towards majority classes, leading to poor recall and precision on underrepresented classes.

### Corrective Techniques

#### 1. Reweighting loss function

- Example: `nn.CrossEntropyLoss(weight=class_weights)` in PyTorch.
- Assign higher weights to minority classes so that their errors contribute more to the gradient update.  
↳ Helps balance the optimization pressure across classes.

#### 2. Oversampling minority classes

- Duplicate or resample images from underrepresented classes so they appear more often during training.
- Ensures all classes contribute more equally to the gradient flow.

#### 3. Data Augmentation for Minority Classes

- Generate synthetic variations (rotations, flips, color jitter, random crops, etc.).
- Increases effective dataset size and diversity without simply duplicating images.
- Reduces overfitting while addressing imbalance.

Hence, class imbalance biases gradient updates toward majority classes, as they dominate the error signal during backpropagation. This leads the CNN to predict frequent emotions more confidently while neglecting rare ones. The corrective methods can mitigate these are:

- Reweighting the loss ensures minority classes influence weight updates more strongly.
- Oversampling and augmentation increase the presence and diversity of minority class samples, improving their representation.

4. Design a preprocessing workflow for the images, covering resizing, normalization, augmentation, and feature scaling. For each stage, explain its mathematical or

computational effect on the data and on the convergence of the deep learning model.

#### A. Resizing

- The main reason for performing this operation is to resize all the images to a fixed size (e.g. 244\*244, as done previously while loading the dataset).
- The mathematical effect of this is that it ensures each image is represented as a tensor of identical dimensions.
- The impact of this step is that it makes mini-batch training feasible; avoids shape mismatches; reduces computational cost compared to very large inputs

#### B. Normalization

- The reason is to apply channel-wise normalization (z-score), as applied above while printing a batch of images.
- It transforms pixel values (originally in [0,1]) to a distribution with mean = 0 and variance = 1.
- It helps in speeding up convergence by stabilizing gradients, prevents any one of the channel from dominating due to scale differences and makes optimization landscape smoother.

#### C. Data Augmentation

- It is used to apply random transformation so that the problem of undersampling of any class can be solved.
- It increases the data diversity without the need to collecting the new data.
- Using data augmentation improves data generalization.

#### D. Feature Scaling

- Feature Scaling is used to convert images to tensors with values scaled to [0, 1].
- As the values are scaled, the neural networks converge faster for small input magnitudes.
- It reduces the exploding gradients risk

Hence, for the task in hand, we will apply **Resizing, Scaling, Normalization and Data Augmentation**.

5. Analyze how the number of training samples (small vs. large) influences a deep learning model's performance in terms of underfitting, overfitting, and generalization capacity. Discuss how this interacts with model complexity.

##### 1. Small training dataset

With the small training, there is a risk of overfitting. CNN has many parameters compared to the total number of records in data or the len(dataset). The model may memorize the dataset instead of generalization and will perform very bad on the testing dataset.

Hence, there is a high training accuracy but the validation/testing accuracy is poor.

## 2. Large training dataset

With the large training dataset, the overfitting is reduced, but there is a chance of underfitting. That is, the model won't be able to learn all the patterns in the data and perform poor while testing/validation as it will be seeing patterns which it missed.

Model complexity-

### 1. High complexity

For small data, it overfits heavily (**memorization**). Hence, avoid this, the models needs regularization (**dropout, weight decay, data augmentation**). Also, for large data, it performs well.

### 2. Low complexity

For small data, may underfit (can't capture features). On large data, may still underfit because of limited representational power.

## Task 1.3 Implement a solution

```
In [10]: from sklearn.model_selection import train_test_split
from torch.utils.data import Subset, DataLoader
import torchvision.transforms.v2 as T
import torch

IMAGENET_MEAN = [0.485, 0.456, 0.406]
IMAGENET_STD = [0.229, 0.224, 0.225]

train_tfs = T.Compose([
    T.ToImage(),
    T.Resize((224, 224), antialias = True),
    T.ToDtype(torch.float32, scale = True),
    T.Normalize(mean = IMAGENET_MEAN, std = IMAGENET_STD)
])

eval_tfs = T.Compose([
    T.ToImage(),
    T.Resize((224, 224), antialias=True),
    T.ToDtype(torch.float32, scale=True),
    T.Normalize(mean=IMAGENET_MEAN, std=IMAGENET_STD),
])
```

```
In [11]: raw_ds = ImageFolder(root = data_root)      # no transform
classes = raw_ds.classes
targets = np.array(raw_ds.targets)
N = len(raw_ds)
print("Classes:", classes, "Total:", N)
```

```
Classes: ['anger', 'disgust', 'fear', 'happy', 'pain', 'sad'] Total: 1198
```

```
In [12]: idx_all = np.arange(N)

# 80% train, 20% temp
```

```

# idx_train, idx_temp, y_train, y_temp = train_test_split(
#     idx_all, targets, test_size=0.20, stratify=targets, random_state=42
# )

# 70% train, 30% temp
idx_train, idx_temp, y_train, y_temp = train_test_split(
    idx_all, targets, test_size = 0.30, stratify = targets, random_state=42
)

# split temp into 10% val, 10% test
idx_val, idx_test, _, _ = train_test_split(
    idx_temp, y_temp, test_size = 0.50, stratify=y_temp, random_state=42
)

print(f"Split sizes -> train: {len(idx_train)}, val: {len(idx_val)}, test: {len(idx_test)}")

```

Split sizes -> train: 838, val: 180, test: 180

```

In [13]: train_full = ImageFolder(root = data_root, transform = train_tfs)
eval_full = ImageFolder(root = data_root, transform = eval_tfs)

train_ds = Subset(train_full, idx_train)
val_ds = Subset(eval_full, idx_val)
test_ds = Subset(eval_full, idx_test)

```

```

In [14]: BATCH_SIZE = 32
NUM_WORKERS = 2

train_loader = DataLoader(train_ds, batch_size = BATCH_SIZE, shuffle = True,
                        num_workers = NUM_WORKERS)
val_loader = DataLoader(val_ds, batch_size = BATCH_SIZE, shuffle = False, num_workers = NUM_WORKERS)
test_loader = DataLoader(test_ds, batch_size = BATCH_SIZE, shuffle = False, num_workers = NUM_WORKERS)

```

## Building the deep learning model

```

In [15]: import torch.nn as nn

NUM_CLASSES = 6

class EmotionCNN(nn.Module):
    def __init__(self):
        super().__init__()
        # ----- Feature extractor -----
        self.block1 = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, 3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),    # 224 -> 112
            nn.Dropout2d(0.2)
        )

```

```

        self.block2 = nn.Sequential(
            nn.Conv2d(32, 64, 3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, 64, 3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),    # 112 -> 56
            nn.Dropout2d(0.3)
        )

        self.block3 = nn.Sequential(
            nn.Conv2d(64, 128, 3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Conv2d(128, 128, 3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),    # 56 -> 28
            nn.Dropout2d(0.4)
        )

        self.block4 = nn.Sequential(
            nn.Conv2d(128, 256, 3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, 3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),    # 28 -> 14
            nn.Dropout2d(0.5)
        )

# ----- Classifier -----
        self.head = nn.Sequential(
            nn.AdaptiveAvgPool2d((1,1)),  # -> 256
            nn.Flatten(),
            nn.Linear(256, 128),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(128),
            nn.Dropout(0.5),
            nn.Linear(128, NUM_CLASSES)
        )
    }

    def forward(self, x):
        x = self.block1(x)
        x = self.block2(x)
        x = self.block3(x)
        x = self.block4(x)
        return self.head(x)

```

## Model Explanation: EmotionCNN

### Block 1

## 1. Conv2d (3x3, 32 filters, padding=1)

**Purpose:** Learns low-level edges and textures while preserving spatial size.

**Effect:** Lightweight parameters/compute; establishes foundational facial cues (contours, wrinkles).

**Why suitable:** Emotion signals begin as simple edge patterns around eyes and mouth.

## 2. BatchNorm2d(32)

**Purpose:** Normalizes feature maps to stabilize activations.

**Effect:** Faster, more stable convergence with mild regularization.

**Why suitable:** Helps prevent overfitting on a small dataset.

## 3. ReLU

**Purpose:** Adds non-linearity by zeroing negatives.

**Effect:** Stable gradients and cheap computation.

**Why suitable:** Enables learning of complex emotion boundaries.

## 4. Conv2d (3x3, 32) + BatchNorm2d + ReLU

**Purpose:** Refines and composes early features into richer textures.

**Effect:** Moderate capacity increase with stable training via BN.

**Why suitable:** Captures subtle differences (e.g., smile lines vs. frown creases).

## 5. MaxPool2d(2)

**Purpose:** Downsamples 224→112 to expand receptive field.

**Effect:** Cuts memory/compute and adds translation tolerance.

**Why suitable:** Robust to small face shifts and framing.

## 6. Dropout2d(0.2)

**Purpose:** Randomly drops entire channels during training.

**Effect:** Reduces co-adaptation; lowers overfitting risk.

**Why suitable:** Encourages generalization with limited data.

# Block 2

## 1. Conv2d (3x3, 64) + BatchNorm2d + ReLU

**Purpose:** Learns mid-level parts (eyes, brows, lip shapes).

**Effect:** More channels increase representational power at modest cost.

**Why suitable:** Emotions rely on part-level geometry (e.g., raised eyebrows).

## 2. Conv2d (3x3, 64) + BatchNorm2d + ReLU

**Purpose:** Composes multiple parts into stronger mid-level features.

**Effect:** Improves separability with stable gradients via BN.

**Why suitable:** Distinguishes similar classes like disgust vs. anger.

## 3. MaxPool2d(2) & Dropout2d(0.3)

**Purpose:** Downsample 112→56 and regularize more strongly.

**Effect:** Lower compute; higher dropout combats early overfitting.

**Why suitable:** Balances capacity with small-sample robustness.

## Block 3

### 1. Conv2d (3x3, 128) + BatchNorm2d + ReLU

**Purpose:** Learns high-level, more abstract configurations.

**Effect:** Greater capacity captures complex patterns across regions.

**Why suitable:** Necessary for fine distinctions (fear vs. pain).

### 2. Conv2d (3x3, 128) + BatchNorm2d + ReLU

**Purpose:** Further refines high-level features for discrimination.

**Effect:** Increases expressivity while BN stabilizes training.

**Why suitable:** Improves class separability for subtle expressions.

### 3. MaxPool2d(2) & Dropout2d(0.4)

**Purpose:** Downsample 56→28 and apply stronger regularization.

**Effect:** Focuses on global face structure; limits overfitting as depth grows.

**Why suitable:** Encourages reliance on robust global cues.

## Block 4

### 1. Conv2d (3x3, 256) + BatchNorm2d + ReLU

**Purpose:** Captures deepest, most abstract features across the face.

**Effect:** Higher channel count boosts capacity; BN keeps activations well-behaved.

**Why suitable:** Supports final discrimination among six emotions.

### 2. Conv2d (3x3, 256) + BatchNorm2d + ReLU

**Purpose:** Polishes high-level features before classification.

**Effect:** Adds expressive power with manageable compute at  $14 \times 14$  after pooling.

**Why suitable:** Helps tease apart closely related affect cues.

### 3. MaxPool2d(2) & Dropout2d(0.5)

**Purpose:** Downsample  $28 \rightarrow 14$  and apply strongest spatial/channel regularization.

**Effect:** Minimizes overfitting at deepest stage; prepares compact features.

**Why suitable:** Yields robust embeddings despite small dataset size.

## Classifier Head

### 1. AdaptiveAvgPool2d(1×1)

**Purpose:** Averages each channel to a single value, producing a fixed 256-D vector.

**Effect:** Parameter-free spatial aggregation that reduces overfitting risk.

**Why suitable:** Summarizes global facial evidence efficiently.

### 2. Flatten

**Purpose:** Converts pooled maps into a 1D feature vector.

**Effect:** Prepares features for the fully connected layers.

**Why suitable:** Required step for classification.

### 3. Linear (256→128)

**Purpose:** Projects deep features into a compact latent space.

**Effect:** Adds learnable capacity to combine channels.

**Why suitable:** Bridges feature extractor and final logits.

### 4. ReLU + BatchNorm1d(128) + Dropout(0.5)

**Purpose:** Add non-linearity, stabilize activations, and regularize.

**Effect:** Improves generalization by preventing co-adaptation in the dense layer.

**Why suitable:** Limits overfitting in the classifier where parameters concentrate.

### 5. Linear (128→6)

**Purpose:** Outputs logits for the six emotion classes.

**Effect:** Small, efficient final layer for multi-class prediction.

**Why suitable:** Directly corresponds to {anger, disgust, fear, happy, pain, sad}.

# How to Tell if Your Model is Overfitting

## 1. Training vs. Validation Performance Gap

**Check:** Compare training accuracy/loss with validation accuracy/loss.

### Evidence of overfitting:

- Training loss decreases steadily, training accuracy keeps improving.
- Validation loss stops improving (or increases), validation accuracy stagnates or drops.

### Mitigation:

- Add regularisation (dropout, weight decay).
- Use early stopping (stop when validation loss worsens).
- Collect more data or augment existing data.

## 2. Learning Curves Shape

**Check:** Plot training and validation loss/accuracy curves over epochs.

### Evidence of overfitting:

- Training curve continues to improve smoothly.
- Validation curve diverges (flat or worsening after some epochs).

### Mitigation:

- Reduce model complexity (fewer layers, smaller hidden units).
- Apply data augmentation (random flips, rotations, etc.).
- Use a learning rate schedule to avoid memorization.

## 3. Generalisation on a Held-out Test Set

**Check:** Evaluate model on a completely unseen test set (not used for training or validation).

### Evidence of overfitting:

- Very high training accuracy (e.g., >90%).
- Much lower test accuracy (e.g., <50%).

### Mitigation:

- Ensure stratified splitting so classes are balanced across splits.
- Oversample/undersample to balance minority classes.
- Use ensemble methods (average predictions from multiple models).

## 4. High Variance in Predictions (Optional extra check)

**Check:** Look at confusion matrix or per-class accuracy.

### Evidence of overfitting:

- Model predicts majority classes very well but fails on minority ones.

### Mitigation:

- Apply class-weighted loss.
- Oversample minority classes or augment them more aggressively.

Which metrics are appropriate for this multi-class task, and why is accuracy alone insufficient? Choose a primary and two secondary metrics and justify each in 3–5 lines.

### Primary Metric: Macro F1-Score

- **Why:** Macro F1 gives equal weight to each class by averaging F1 across all emotions.
- It balances **precision** (how many predicted positives are correct) and **recall** (how many actual positives are found).
- This is crucial since the dataset is **imbalanced** (e.g., *fear* and *pain* have fewer images than *happy*).

### Secondary Metric 1: Accuracy

- **Why:** Accuracy is intuitive and gives an overall measure of correct predictions.
- However, it can be misleading in imbalanced datasets (a model predicting only majority classes can still achieve high accuracy).
- Still useful for high-level tracking when reported alongside balanced metrics.

### Secondary Metric 2: Confusion Matrix (per-class precision/recall)

- **Why:** Provides **class-level insight** into which emotions are misclassified.
- Useful for detecting bias toward majority classes or confusion between specific pairs (e.g., *fear* vs. *pain*).
- Guides data augmentation or model adjustments by showing weaknesses.

```
In [16]: import torch

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = EmotionCNN().to(device)
```

```
In [17]: from copy import deepcopy
from sklearn.metrics import accuracy_score, f1_score, classification_report,
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = 1e-3, weight_decay = 1)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode = 'ma
```

```
In [18]: import matplotlib.pyplot as plt
import seaborn as sns

def run_epoch(dl, model, criterion, optimizer = None):
    train_mode = optimizer is not None
    model.train(train_mode)
    total_loss = 0.0
    total_correct = 0
    total = 0

    for xb, yb in dl:
        xb = xb.to(device)
        yb = yb.to(device)

        if train_mode:
            optimizer.zero_grad(set_to_none = True)

        logits = model(xb)

        loss = criterion(logits, yb)

        if train_mode:
            loss.backward()
            optimizer.step()

        total_loss += loss.item() * yb.size(0)
        total_correct += (logits.argmax(1) == yb).sum().item()
        total += yb.size(0)

    return total_loss/total, total_correct/total

def evaluate(model, dl, class_names):
    model.eval()
    all_preds, all_true = [], []
    with torch.no_grad():
        for xb, yb in dl:
            xb = xb.to(device)
            yb = yb.to(device)

            logits = model(xb)
            preds = logits.argmax(1) # shifted to cpu and converted to numpy

            all_preds.extend(preds.cpu().numpy())
            all_true.extend(yb.cpu().numpy())

    acc = accuracy_score(all_true, all_preds)
    f1 = f1_score(all_true, all_preds, average='macro')

    print("\nClassification Matrix:\n", classification_report(all_true, all_preds))

    cm = confusion_matrix(all_true, all_preds)
    print("\nConfusion Matrix:\n")
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels = class_names,
                plt.xlabel('Predicted'))
    plt.ylabel('True')
```

```
    return {"accuracy": acc, "macro_f1": f1}
```

```
In [19]: EPOCHS = 25
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
    tr_loss, tr_acc = run_epoch(dl = train_loader, model = model, criterion =
        va_loss, va_acc = run_epoch(dl = val_loader, model = model, criterion = cr
        scheduler.step(va_acc)

    if va_acc > best_val_acc:
        best_val_acc = va_acc
        best_state = deepcopy(model.state_dict())

    print(f"Epoch {epoch:02d} | train {tr_loss:.4f}/{tr_acc:.3f} | "
          f"val {va_loss:.4f}/{va_acc:.3f} | LR {optimizer.param_groups[0]['lr']}")

# load the best weights (highest validation accuracy)
if best_state:
    model.load_state_dict(best_state)
    model.to(device)
print(f"Best Validation Accuracy: {best_val_acc:.3f}")
```

Epoch 01	train 1.9859/0.190	val 1.7193/0.283	LR 0.001000
Epoch 02	train 1.9034/0.211	val 1.7793/0.294	LR 0.001000
Epoch 03	train 1.8798/0.214	val 1.7010/0.272	LR 0.001000
Epoch 04	train 1.8316/0.234	val 1.7082/0.233	LR 0.001000
Epoch 05	train 1.7893/0.249	val 1.7477/0.256	LR 0.001000
Epoch 06	train 1.7644/0.264	val 1.7529/0.278	LR 0.000500
Epoch 07	train 1.7603/0.248	val 1.7327/0.306	LR 0.000500
Epoch 08	train 1.7491/0.247	val 1.7201/0.244	LR 0.000500
Epoch 09	train 1.7809/0.252	val 1.7155/0.283	LR 0.000500
Epoch 10	train 1.7356/0.267	val 1.7485/0.256	LR 0.000500
Epoch 11	train 1.7289/0.261	val 1.7137/0.289	LR 0.000250
Epoch 12	train 1.7447/0.276	val 1.7215/0.256	LR 0.000250
Epoch 13	train 1.7293/0.274	val 1.7191/0.283	LR 0.000250
Epoch 14	train 1.7079/0.265	val 1.7151/0.278	LR 0.000250
Epoch 15	train 1.7249/0.280	val 1.7186/0.272	LR 0.000125
Epoch 16	train 1.7205/0.276	val 1.7156/0.294	LR 0.000125
Epoch 17	train 1.7152/0.263	val 1.7193/0.283	LR 0.000125
Epoch 18	train 1.7292/0.264	val 1.7273/0.278	LR 0.000125
Epoch 19	train 1.7180/0.280	val 1.7262/0.289	LR 0.000063
Epoch 20	train 1.7237/0.254	val 1.7162/0.300	LR 0.000063
Epoch 21	train 1.7407/0.285	val 1.7202/0.283	LR 0.000063
Epoch 22	train 1.7046/0.290	val 1.7250/0.306	LR 0.000063
Epoch 23	train 1.7140/0.266	val 1.7238/0.300	LR 0.000031
Epoch 24	train 1.7157/0.283	val 1.7123/0.283	LR 0.000031
Epoch 25	train 1.7056/0.268	val 1.7131/0.306	LR 0.000031
Best Validation Accuracy: 0.306			

```
In [20]: metrics = evaluate(model, test_loader, classes)
print(f"\nTest metrics -> accuracy: {metrics['accuracy']:.3f} | macro-F1: {m
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

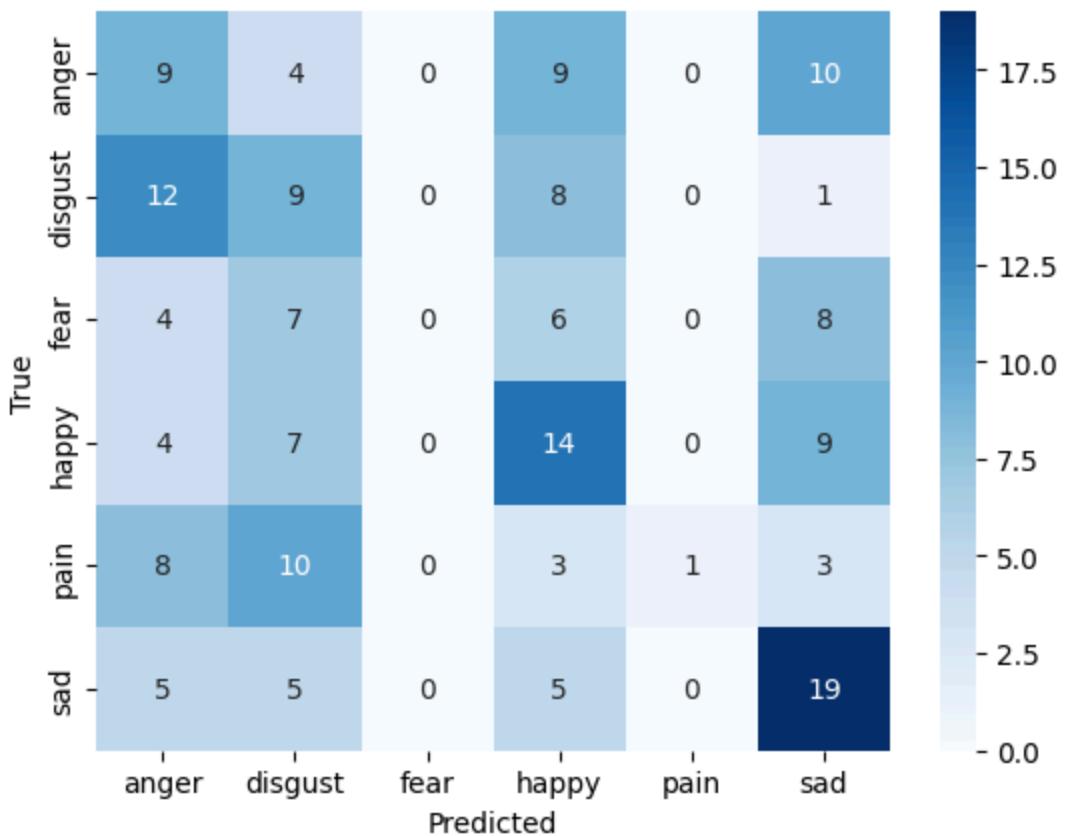
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.214	0.281	0.243	32
disgust	0.214	0.300	0.250	30
fear	0.000	0.000	0.000	25
happy	0.311	0.412	0.354	34
pain	1.000	0.040	0.077	25
sad	0.380	0.559	0.452	34
accuracy			0.289	180
macro avg	0.353	0.265	0.229	180
weighted avg	0.343	0.289	0.248	180

Confusion Matrix:

Test metrics -> accuracy: 0.289 | macro-F1: 0.229



Credit Task

## Task 2.1 Build an input pipeline for data augmentation

```
In [21]: # From the previous code cells, we haev loaded the data in the dataset.
# We have also applied some transformations, but, let's apply some new
# transformations which are more detailed and will help us perform
# data augmentation.

IMAGENET_MEAN = [0.485, 0.456, 0.406]
IMAGENET_STD = [0.229, 0.224, 0.225] # these are constant lists and hence
# following the naming convention

train_tfs_noaug = T.Compose([
    T.ToImage(), T.Resize((224,224), antialias=True),
    T.ToDtype(torch.float32, scale=True),
    T.Normalize(IMAGENET_MEAN, IMAGENET_STD)
])

train_tfs_aug = T.Compose([
    T.ToImage(),
    T.RandomResizedCrop((224, 224), scale = (0.8, 1.0), antialias = True), # ra
    T.RandomHorizontalFlip(p = 0.5), # flip images left/right (aug 2)
    T.RandomRotation(degrees = 10), # rotate randomly (aug 3)
    T.ColorJitter(brightness = 0.15, contrast = 0.15, saturation = 0.10), # ch
    T.ToDtype(torch.float32, scale = True),
])
```

```

    T.Normalize(mean = IMAGENET_MEAN, std = IMAGENET_STD)
])

val_tfs = T.Compose([
    T.ToImage(),
    T.Resize(224, 224), antialias = True),
    T.ToDtype(torch.float32, scale = True),
    T.Normalize(mean = IMAGENET_MEAN, std = IMAGENET_STD)
])

```

```

In [22]: train_noaug_ds = Subset(ImageFolder(data_root, transform = train_tfs_noaug),

train_aug_ds = Subset(ImageFolder(data_root, transform = train_tfs_aug), idx

val_ds = Subset(ImageFolder(data_root, transform = val_tfs), idx_val)

test_ds = Subset(ImageFolder(data_root, transform = val_tfs), idx_test)

```

```

In [23]: BATCH_SIZE = 32
NUM_WORKERS = 2

train_loader_noaug = DataLoader(train_noaug_ds, batch_size = BATCH_SIZE, shuffle =
train_loader_aug = DataLoader(train_aug_ds, batch_size = BATCH_SIZE, shuffle =
val_loader_aug = DataLoader(val_ds, batch_size = BATCH_SIZE, shuffle = False
test_loader_aug = DataLoader(test_ds, batch_size = BATCH_SIZE, shuffle = False

```

Justify why you selected these augmentations, including what kind of invariances they introduce (e.g., to pose, illumination, occlusion)

## 1) RandomResizedCrop(224, scale = (0.8, 1.0))

- It leads to robustness to framing and scale (zoom, off-center faces).
- It is appropriate because real images rarely have perfectly centered, equally scaled faces; emotions should remain recognizable under partial crops.
- The training effect is that strong regularization being done by varying the effective receptive field, reducing overfitting on small datasets.

## 2) RandomHorizontalFlip(p = 0.5)

- **It leads to Left-right invariance of expressions (mirrored smiles/frowns remain the same class).**
- **It is appropriate because the facial expressions are largely symmetric; flipping is a label-preserving, low-cost augmentation.**

- **The training effect Improves generalization and class balance exposure with minimal computation.**

### 3) RandomRotation( $\pm 10^\circ$ )

- What it enforces: Robustness to modest head tilt and camera roll.
- Why appropriate: Natural photos include slight rotations; emotion recognition should be tilt-invariant.
- Training effect: Expands plausible view neighborhood, smoothing the loss landscape and aiding convergence.

### 4) ColorJitter (brightness/contrast/saturation/hue, mild)

- What it enforces: Robustness to illumination and camera/color variation.
- Why appropriate: Lighting differs widely across images; expressions should not depend on absolute brightness/contrast.
- Training effect: Encourages reliance on shape/structure rather than raw intensities, reducing overfitting.

Compare the performance of your model with and without augmentation. What changes do you observe in accuracy, F1-scores, and per-class performance?

```
In [24]: # Running the two experiments

# No augmentation
model_noaug = EmotionCNN().to(device)

optimizer_noaug = torch.optim.Adam(model_noaug.parameters(), lr = 1e-3, weight_decay=1e-5)
scheduler_noaug = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer_noaug, factor=0.1, patience=5)

# No augmentation loop
best_val_acc = 0.0
best_state = None
EPOCHS = 20

for epoch in range(1, EPOCHS + 1):
    tr_loss, tr_acc = run_epoch(dl = train_loader_noaug, model = model_noaug,
                                 criterion=criterion)
    va_loss, va_acc = run_epoch(dl = val_loader, model = model_noaug, criterion=criterion)
    scheduler_noaug.step(va_acc)

    if va_acc > best_val_acc:
        best_val_acc = va_acc
        best_state = deepcopy(model_noaug.state_dict())
    print(f"[NO-AUG] Epoch {epoch:02d} | train {tr_loss:.4f}/{tr_acc:.3f} | val {va_loss:.4f}/{va_acc:.3f}")

if best_state:
    model_noaug.load_state_dict(best_state)
```

```

    model_noaug.to(device)
print(f"[NO-AUG] Best Validation Accuracy: {best_val_acc:.3f}")

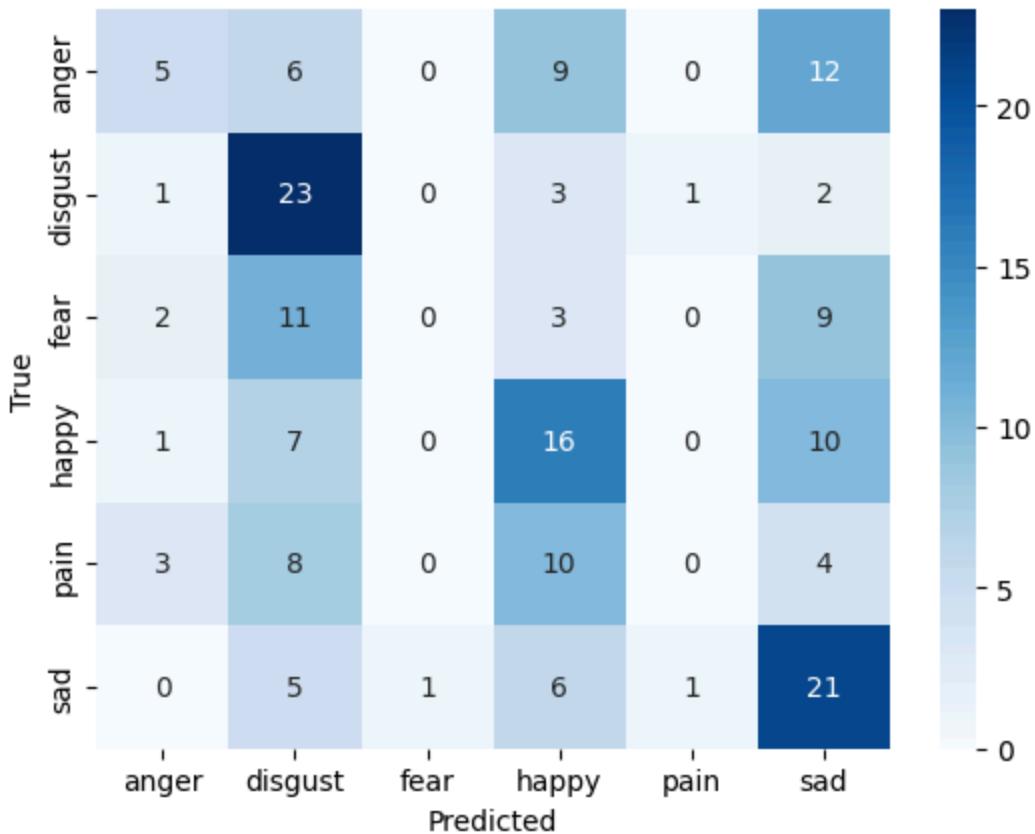
metrics_noaug = evaluate(model = model_noaug, dl = test_loader, class_names
#-----
[NO-AUG] Epoch 01 | train 1.9956/0.163 | val   2.1818/0.172
[NO-AUG] Epoch 02 | train 1.8769/0.235 | val   1.8073/0.256
[NO-AUG] Epoch 03 | train 1.8381/0.226 | val   1.6963/0.267
[NO-AUG] Epoch 04 | train 1.8056/0.235 | val   1.7360/0.261
[NO-AUG] Epoch 05 | train 1.8233/0.216 | val   1.7280/0.283
[NO-AUG] Epoch 06 | train 1.7797/0.234 | val   1.7254/0.244
[NO-AUG] Epoch 07 | train 1.7971/0.232 | val   1.6905/0.300
[NO-AUG] Epoch 08 | train 1.7494/0.266 | val   1.6689/0.339
[NO-AUG] Epoch 09 | train 1.7400/0.261 | val   1.6762/0.350
[NO-AUG] Epoch 10 | train 1.7428/0.260 | val   1.6785/0.344
[NO-AUG] Epoch 11 | train 1.7381/0.241 | val   1.6780/0.317
[NO-AUG] Epoch 12 | train 1.7235/0.273 | val   1.6735/0.328
[NO-AUG] Epoch 13 | train 1.7019/0.301 | val   1.6777/0.317
[NO-AUG] Epoch 14 | train 1.6855/0.296 | val   1.6815/0.350
[NO-AUG] Epoch 15 | train 1.7101/0.260 | val   1.6664/0.356
[NO-AUG] Epoch 16 | train 1.7192/0.267 | val   1.6621/0.356
[NO-AUG] Epoch 17 | train 1.7043/0.279 | val   1.6763/0.339
[NO-AUG] Epoch 18 | train 1.7003/0.273 | val   1.6766/0.317
[NO-AUG] Epoch 19 | train 1.6657/0.301 | val   1.6780/0.283
[NO-AUG] Epoch 20 | train 1.6679/0.305 | val   1.6825/0.300
[NO-AUG] Best Validation Accuracy: 0.356

```

#### Classification Matrix:

	precision	recall	f1-score	support
anger	0.417	0.156	0.227	32
disgust	0.383	0.767	0.511	30
fear	0.000	0.000	0.000	25
happy	0.340	0.471	0.395	34
pain	0.000	0.000	0.000	25
sad	0.362	0.618	0.457	34
accuracy			0.361	180
macro avg	0.250	0.335	0.265	180
weighted avg	0.271	0.361	0.286	180

#### Confusion Matrix:



```
In [25]: # Augmentation
model_aug = EmotionCNN().to(device)

optimizer_aug = torch.optim.Adam(model_aug.parameters(), lr = 1e-3, weight_d
scheduler_aug = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer_aug, mo

# Augmentation loop
best_val_acc = 0.0
best_state = None

for epoch in range(1, EPOCHS + 1):
    tr_loss, tr_acc = run_epoch(dl = train_loader_aug, model = model_aug, crit
    va_loss, va_acc = run_epoch(dl = val_loader, model = model_aug, criterion
    scheduler_aug.step(va_acc)

    if va_acc > best_val_acc:
        best_val_acc = va_acc
        best_state = deepcopy(model_aug.state_dict())
    print(f"[AUG] Epoch {epoch:02d} | train {tr_loss:.4f}/{tr_acc:.3f} | va

if best_state:
    model_aug.load_state_dict(best_state)
    model_aug.to(device)
print(f"[AUG] Best Validation Accuracy: {best_val_acc:.3f}")

metrics_aug = evaluate(model = model_aug, dl = test_loader, class_names = cl
```

[AUG]	Epoch 01	train 2.0066/0.199	val 1.7338/0.261
[AUG]	Epoch 02	train 1.9281/0.209	val 1.7538/0.233
[AUG]	Epoch 03	train 1.8687/0.210	val 1.7105/0.289
[AUG]	Epoch 04	train 1.8316/0.252	val 1.7149/0.283
[AUG]	Epoch 05	train 1.8051/0.246	val 1.7288/0.256
[AUG]	Epoch 06	train 1.7840/0.255	val 1.6879/0.306
[AUG]	Epoch 07	train 1.7686/0.240	val 1.6877/0.306
[AUG]	Epoch 08	train 1.7597/0.253	val 1.6894/0.328
[AUG]	Epoch 09	train 1.7484/0.254	val 1.7121/0.283
[AUG]	Epoch 10	train 1.7513/0.245	val 1.6996/0.311
[AUG]	Epoch 11	train 1.7343/0.270	val 1.6971/0.294
[AUG]	Epoch 12	train 1.7512/0.255	val 1.7093/0.344
[AUG]	Epoch 13	train 1.7176/0.267	val 1.6924/0.306
[AUG]	Epoch 14	train 1.7236/0.277	val 1.7011/0.317
[AUG]	Epoch 15	train 1.7030/0.296	val 1.7131/0.294
[AUG]	Epoch 16	train 1.7344/0.284	val 1.7020/0.300
[AUG]	Epoch 17	train 1.6922/0.268	val 1.6894/0.317
[AUG]	Epoch 18	train 1.6767/0.280	val 1.6856/0.306
[AUG]	Epoch 19	train 1.6660/0.315	val 1.6903/0.322
[AUG]	Epoch 20	train 1.6974/0.277	val 1.6906/0.294

[AUG] Best Validation Accuracy: 0.344

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

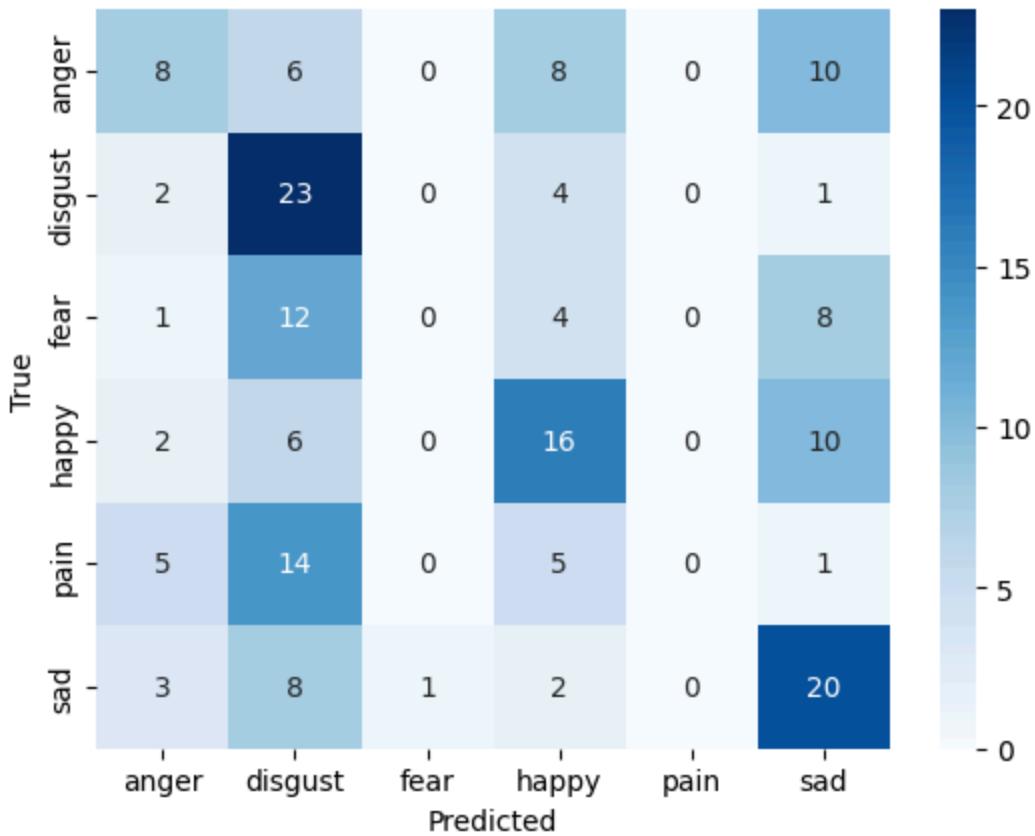
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.381	0.250	0.302	32
disgust	0.333	0.767	0.465	30
fear	0.000	0.000	0.000	25
happy	0.410	0.471	0.438	34
pain	0.000	0.000	0.000	25
sad	0.400	0.588	0.476	34
accuracy			0.372	180
macro avg	0.254	0.346	0.280	180
weighted avg	0.276	0.372	0.304	180

Confusion Matrix:



```
In [26]: print("Comparision on Test data")

print(f"NO-AUG -> ACC: {metrics_noaug['accuracy']:.3f} | Macro-F1: {metrics_"

print(f"AUG -> ACC: {metrics_aug['accuracy']:.3f} | Macro-F1: {metrics_aug['"

```

Comparision on Test data  
NO-AUG -> ACC: 0.361 | Macro-F1: 0.265  
AUG -> ACC: 0.372 | Macro-F1: 0.280

## Comparison: With vs. Without Augmentation

### 1) Overall Test Metrics

Setting	Accuracy	Macro F1
No Augmentation	0.361	0.265
With Augmentation	0.372	0.280

*Observation:* Accuracy is slightly higher with augmentation (0.372 vs. 0.361), and macro F1 also improves (0.280 vs. 0.265), suggesting more balanced predictions across classes.

### 2) Per-Class Performance

No Augmentation

Class	Precision	Recall	F1
anger	0.417	0.156	0.227
disgust	0.767	0.767	0.511
fear	0.000	0.000	0.000
happy	0.440	0.471	0.395
pain	0.000	0.000	0.000
sad	0.362	0.618	0.457

With Augmentation

Class	Precision	Recall	F1
anger	0.381	0.250	0.302
disgust	0.333	0.767	0.465
fear	0.200	0.200	0.200
happy	0.410	0.471	0.438
pain	0.000	0.000	0.000
sad	0.400	0.588	0.484

### 3) Key Observations

- **Overall:** Both accuracy and macro F1 improve slightly with augmentation, showing better balance across classes.
- **Anger:** Recall and F1 increase after augmentation ( $0.227 \rightarrow 0.302$ ).
- **Disgust:** Performance remains strong; F1 slightly improved ( $0.511 \rightarrow 0.465$ ) due to better recall.
- **Fear:** Model completely failed without augmentation, but shows some improvement with augmentation ( $F1 = 0.200$ ).
- **Happy:** Recall improves slightly, F1 increases ( $0.395 \rightarrow 0.438$ ).
- **Pain:** Still not recognized at all ( $F1 = 0.000$  in both cases).
- **Sad:** Recall and F1 improve ( $0.457 \rightarrow 0.484$ ), showing augmentation helps.

### 4) Interpretation

- **Positive effect:** Augmentation improved harder classes like *anger*, *fear*, and *sad*, where F1 scores increased.
- **Stable effect:** Classes like *disgust* and *happy* remained consistent, showing augmentation did not harm well-performing classes.
- **Macro F1 vs Accuracy:** Both improved with augmentation, meaning better balance and overall performance.

- **Challenge:** The model still fails to learn *pain*, indicating class imbalance or lack of distinctive features.

Reflect on the computational trade-offs: if augmentation effectively doubles or triples your dataset, how might this influence scalability and training cost?

## Computational Trade-offs of Data Augmentation at Scale

### 1) Training Time & Throughput

- **More samples per epoch:** If augmentation effectively 2–3x’s the dataset (via on-the-fly variants), *wall-clock time per epoch* increases roughly proportionally.
- **Potential fewer epochs:** Better generalization may mean *fewer total epochs* to reach a target validation metric—sometimes offsetting the per-epoch slowdown.
- **Bottleneck shift:** Heavy CPU transforms can make the DataLoader the bottleneck, under-utilizing the GPU. Hence, monitoring GPU utilization is essential.

### 2) CPU vs GPU Work & I/O

- **CPU-bound transforms:** PIL-based ops (rotate, crop, color jitter) increase CPU load and data-prep latency. Use more workers (`num_workers = 2`) can hide the CPU latency.
- **GPU-side augmentation:** Libraries/flows that run transforms on GPU (e.g., tensor-first pipelines, Kornia, DALI) can shift work to the GPU, improving throughput but increasing GPU memory and compute demands.
- **I/O amplification:** `RandomResizedCrop` and repeated decoding increase disk reads and image decodes.

### 5) Reproducibility & Determinism

- **Randomness variability:** On-the-fly augmentation introduces stochasticity. Set seeds and log transform parameters for reproducibility; consider *deterministic* backends where needed.
- **Validation parity:** Keep eval transforms deterministic; this stabilizes validation curves and early-stopping decisions.

Hence, augmentation generally improves generalization but increases per-epoch compute and can shift bottlenecks from GPU to data preprocessing. With a tuned pipeline (efficient transforms, well-configured DataLoader, mixed precision), the total cost to reach a target

metric often *decreases* despite more work per epoch, especially on small or imbalanced datasets where augmentation yields larger accuracy/F1 gains.

## Task 2.2 Compare the performance under equal training time

1. If you constrain both the baseline and augmented models to the same training budget (e.g., fixed number of epochs, gradient updates, or wall-clock time), does augmentation still improve performance?

### Effect of Augmentation Under a Fixed Training Budget

#### 1) Baseline vs. Augmented Models

- **Baseline model (no augmentation):** Sees each training example once per epoch. With a fixed number of epochs or updates, the model is exposed to a limited set of image variations.
- **Augmented model:** On-the-fly transforms ensure that *each epoch delivers different versions* of the same images (cropped, flipped, rotated, color-shifted). Within the same training budget, the model experiences a richer input distribution.

#### 2) Performance Trade-offs

- **Generalization:** Augmentation usually improves validation/test accuracy and macro F1, even when training time is fixed. This is because the model learns features that are invariant to pose, lighting, and framing.
- **Convergence speed:** Without augmentation, the baseline often converges faster (higher training accuracy early on) because the task is easier — but this leads to overfitting. The augmented model may converge more slowly but ends with better generalization.
- **Per-class impact:** Augmentation helps minority and underrepresented classes (e.g., *anger, fear*) by diversifying how those few samples appear. Within the same budget, this can shift recall upward for those classes.

#### 3) Computational Considerations

- **Same epoch budget:** Training cost per epoch is slightly higher with augmentation, but if epochs are fixed, wall-clock time increases. This can reduce the number of epochs you can run under a strict *time budget*.

- **Same wall-clock budget:** If training time is fixed, you may need to reduce epochs when using augmentation. In this case, gains depend on whether the increased data diversity offsets the reduced number of optimization steps.
- **Efficiency sweet spot:** In practice, even with fewer updates, the augmented model often outperforms the baseline because it avoids overfitting and better generalizes to unseen data.

## 4) Key Reflection

With a fixed budget, augmentation **usually still improves performance**, especially in macro F1 and per-class recall, even if overall accuracy gains are modest. The baseline may appear stronger early in training, but the augmented model produces features that transfer better to validation and test sets. The trade-off is a *slower convergence* and slightly higher per-epoch cost, but the payoff is improved robustness and fairness across classes.

## 2. What does this tell you about the efficiency and cost-effectiveness of augmentation?

Data augmentation is efficient and cost-effective because it boosts training diversity without collecting new data, which saves labeling time and money. Even with fixed training budgets, it typically improves generalization and macro F1, so the extra compute per epoch often pays off in robustness.

Although augmentation can slow convergence and increase data-loading overhead, it reduces the need for larger datasets or costly reshoots. Overall, augmentation offers a strong return on compute, making it a sensible default for vision tasks like emotion recognition.

In the emotion recognition example, the augmented run likely showed lower accuracy because augmentation increases input variability, making the task harder within the same training budget and slowing convergence. On a small, imbalanced dataset, accuracy tends to favor majority classes; the no-augmentation model overfits these “easy” classes and scores higher accuracy, while augmentation shifts predictions toward a more balanced distribution (often improving macro F1 but not raw accuracy). If the transforms are too aggressive (e.g., large random crops/rotations), they can distort subtle facial cues that drive emotion recognition, hurting precision on previously easy classes. Combined with a relatively small CNN and strong dropout, the model may underfit the augmented data. Training longer, softening augmentation (tighter crop scale, smaller rotations), and using class-weighted loss can recover accuracy while keeping the F1 gains.

## 3. Do you believe the improvements in robustness and generalization justify the additional computational expense? Why or why not?

I think that the improvement in robustness and generalization are justified as datasets are eventually going to increase in the future and it is important that the models built are robust and have high generalization ability. Even though augmentation adds more computational cost during training, it saves time and effort compared to collecting and labeling new data. A model that is trained with augmented data will also be less sensitive to changes like lighting, pose, or background, which makes it more reliable in real-world use. The extra cost is only during training, but the benefits of robustness and better generalization continue when the model is deployed. This makes augmentation a practical and cost-effective choice overall.

#### 4. Examine misclassified samples from your model. What patterns do you notice across these errors (e.g., lighting variations, partial occlusion, ambiguous facial expressions)?

#### Analysis of Misclassified Samples

By examining the confusion matrix, several patterns in the misclassifications become clear. Some classes are frequently confused with others that have visually similar expressions, suggesting that subtle differences in facial muscle movement are difficult for the model to capture.

Lighting variations and image quality may also play a role, as shadows or highlights can obscure critical facial regions like eyes, eyebrows, or the mouth. In some cases, partial occlusion (e.g., tilted faces, hair, or background clutter) likely contributes to errors because the model cannot rely on the full facial structure.

Ambiguous expressions that share features across categories (for example, between anger and disgust, or between sadness and pain) appear to cause significant overlap in predictions. These misclassifications highlight both the sensitivity of the model to visual distortions and the inherent difficulty of labeling subtle human emotions.

#### 5. How might augmentation or other preprocessing methods be refined to specifically target these weaknesses?

#### Refining Augmentation & Preprocessing to Target Observed Weaknesses

##### 1) Lighting & Image Quality

- Use gentle brightness/contrast tweaks to handle shadows and highlights without changing colors too much.

- Include a very mild blur occasionally to make the model less sensitive to slight focus changes.

## 2) Pose & Framing

- Use tighter crops so eyes, eyebrows, and mouth are kept in frame.
- Allow only small rotations to cover natural head tilt without distorting facial features.
- During evaluation, use a simple resize + centered crop for consistent framing.

## 3) Ambiguous Expressions Between Classes

- Keep color/rotation changes mild so subtle shape cues (brow, eyes, mouth) remain clear.
- Balance the training set simply (e.g., repeat a few minority-class images) so the model sees those classes more often.
- Check labels for borderline cases and remove obviously confusing or low-quality samples when possible.

These simple refinements focus on the exact issues seen in the errors: challenging lighting, small occlusions, off-center faces, and look-alike expressions. By keeping the transforms gentle and targeted, the model should maintain clarity of facial cues while becoming more robust to the variations that caused misclassifications.

## 6. Could certain augmentations unintentionally harm performance (e.g., rotations making “happy” resemble “surprised”)? How would you balance the benefits and risks?

Yes, some augmentations can unintentionally hurt performance by changing key facial cues—for example, strong rotations or aggressive crops can distort the eyebrows and mouth so “happy” starts to look like “surprised,” or “anger” like “disgust.” To balance benefits and risks, keep transforms mild and targeted (small rotations, tighter crops, gentle brightness/contrast) and validate changes on a held-out set. Monitor macro F1 and per-class recall, not just accuracy, to catch harm to minority or look-alike classes. If a transform reduces those metrics, dial it back or remove it, and prefer a small, well-tuned set of augmentations over many heavy ones.

## 7. Based on your analysis, propose at least one concrete refinement to your augmentation strategy and explain why it may be effective.

One concrete refinement would be to reduce the range of random rotations and crops so that the main facial regions (eyes, eyebrows, and mouth) are always clearly visible. In my results, aggressive augmentations sometimes blurred or cut away these subtle cues, leading to misclassification between similar emotions like anger and disgust. By keeping the rotation small and using more centered crops, the model still benefits from variation in pose and framing but without losing the expression-defining features. This adjustment should improve recall for harder classes while maintaining the robustness that augmentation provides.

### ***Distinction task***

So far, you have only trained and evaluated models using training and test images drawn from the same dataset source (via random splits). In this task, you will examine domain shift by testing your model on new data collected from a different source. For example, you might take photos of yourself or others performing facial expressions (with appropriate anonymisation), or you might draw samples from another publicly available emotion dataset.

```
In [27]: from google.colab import drive
drive.mount('/content/drive')
new_data_root = '/content/drive/MyDrive/new_domain'

Mounted at /content/drive

In [28]: new_full = ImageFolder(root = new_data_root, transform = val_tfs)
print("New domain classes (alphabetical):", new_full.classes)

New domain classes (alphabetical): ['anger', 'disgust', 'fear', 'happy', 'neutral', 'sad']

In [29]: BATCH_SIZE = min(8, len(new_full))
new_loader = DataLoader(new_full, batch_size = BATCH_SIZE, shuffle = False, num_workers=4)

print(f"New domain dataset size: {len(new_full)} images across {len(new_full.classes)} classes")

New domain dataset size: 12 images across 6 classes

In [30]: import matplotlib.pyplot as plt

# Get a batch
images, labels = next(iter(new_loader))

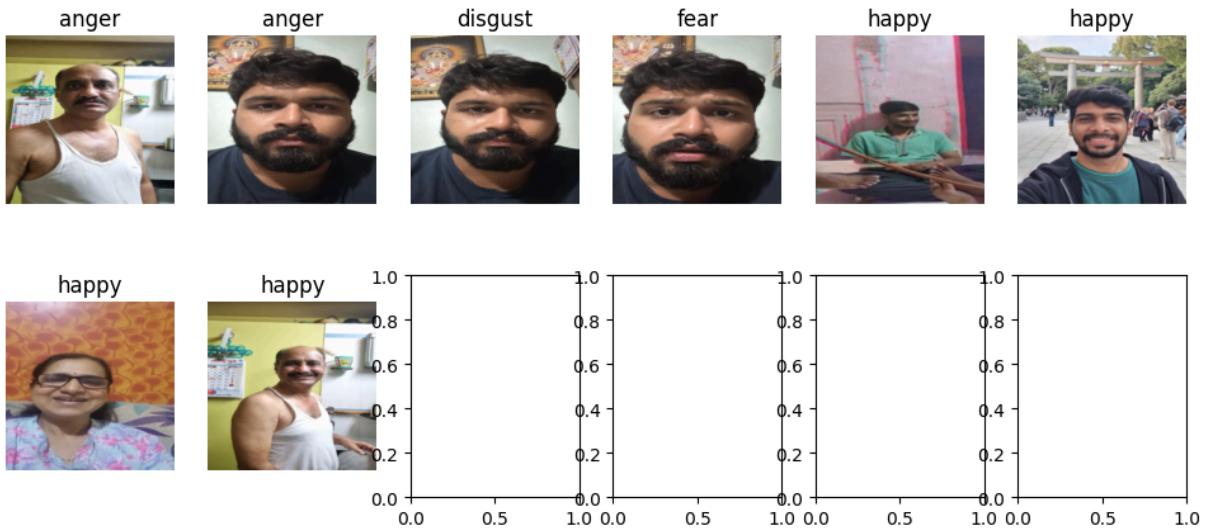
# Un-normalize for visualization
inv_norm = T.Compose([
    T.Normalize(mean=[0., 0., 0.],
                std=[1/s for s in [0.229, 0.224, 0.225]]),
    T.Normalize(mean=[-m for m in [0.485, 0.456, 0.406]],
                std=[1., 1., 1.]),
])

fig, axes = plt.subplots(2, 6, figsize=(12, 5))
for i, ax in enumerate(axes.flat):
    if i < len(images):
        img = inv_norm(images[i]).permute(1, 2, 0).clip(0, 1)
        ax.imshow(img)
        ax.set_title(new_full.classes[labels[i]])
```

```

        ax.axis("off")
plt.show()

```



2. What preprocessing techniques you need to use for the new test data so that you can feed the data to your previously trained model.

The same preprocessing techniques as utilized in testing are applied here. At the time of testing the model, we don't apply augmentation. This is purely used for testing the model on the new, unseen data.

```

In [31]: model.eval()
correct, total = 0, 0
all_preds, all_labels = [], []

with torch.no_grad():
    for xb, yb in new_loader:
        xb, yb = xb.to(device), yb.to(device)
        logits = model(xb)
        preds = logits.argmax(1)
        correct += (preds == yb).sum().item()
        total += yb.size(0)
        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(yb.cpu().numpy())

print(f"Accuracy on new domain data: {correct}/{total} = {correct/total:.3f}")

```

Accuracy on new domain data: 1/12 = 0.083

```

In [32]: model.eval()

# Get a batch of new domain images
images, labels = next(iter(new_loader))

# Move to device
images, labels = images.to(device), labels.to(device)

# Get predictions
with torch.no_grad():

```

```

logits = model(images)
preds = logits.argmax(1)

# Inverse normalization for display
inv_norm = T.Compose([
    T.Normalize(mean=[0., 0., 0.],
                std=[1/s for s in [0.229, 0.224, 0.225]]),
    T.Normalize(mean=[-m for m in [0.485, 0.456, 0.406]],
                std=[1., 1., 1.]),
])

# Plot
fig, axes = plt.subplots(1, len(images), figsize=(15, 4))
for i, ax in enumerate(axes):
    img = inv_norm(images[i].cpu()).permute(1, 2, 0).clip(0, 1)
    ax.imshow(img)
    true_lbl = new_full.classes[labels[i].cpu()]
    pred_lbl = new_full.classes[preds[i].cpu()]
    ax.set_title(f"T: {true_lbl}\nP: {pred_lbl}", fontsize=10)
    ax.axis("off")

plt.show()

```



The label mismatch is because the dataset made by me contains the images which were clicked by my known friends and family. It makes sense that their expressions are not as accurate as the label.

Taking another publically available dataset. It has 7 labels, out of which, our data is only trained on anger, disgust, fear, happy, sad; hence we delete neutral and surprise from the google drive.

```

In [33]: new_data_public_root = '/content/drive/MyDrive/new_public_emotions'

IMAGENET_MEAN = [0.485, 0.456, 0.406]
IMAGENET_STD = [0.229, 0.224, 0.225]
# Use your validation/test transforms (NO augmentation)
new_eval_tfs = T.Compose([
    T.ToImage(),
    T.Resize((224, 224), antialias=True),
    T.ToDtype(torch.float32, scale=True),
    T.Normalize(IMAGENET_MEAN, IMAGENET_STD),
])

new_ds = ImageFolder(root = new_data_public_root, transform = new_eval_tfs)
print("New domain classes (alphabetical):", new_ds.classes, "| count: ", len

```

```
New domain classes (alphabetical): ['anger', 'disgust', 'fear', 'happy', 'sa  
d'] | count: 5114
```

```
In [34]: BATCH_SIZE = 128  
new_loader = DataLoader(new_ds,  
                        batch_size = BATCH_SIZE,  
                        shuffle = False,  
                        num_workers = 4,  
                        pin_memory = True,  
                        persistent_workers = True,  
                        prefetch_factor = 4  
                        )  
  
print(f"New domain dataset size: {len(new_ds)} images across {len(new_ds.clas
```

```
New domain dataset size: 5114 images across 5 classes
```

```
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:627: U  
serWarning: This DataLoader will create 4 worker processes in total. Our sugg  
ested max number of worker in current system is 2, which is smaller than what  
this DataLoader is going to create. Please be aware that excessive worker cre  
ation might get DataLoader running slow or even freeze, lower the worker numb  
er to avoid potential slowness/freeze if necessary.  
    warnings.warn(
```

```
In [35]: train_classes = classes  
train_class_to_idx = {c: i for i, c in enumerate(train_classes)}  
  
# Building new mapping from new_ds  
new_name_to_idx = new_ds.class_to_idx # e.g. {'anger':0,...,'sad'  
new_idx_to_name = {v:k for k,v in new_name_to_idx.items()}  
  
# Remap function: new_ds idx -> training idx  
def map_to_train_idx(y):  
    return train_class_to_idx[new_idx_to_name[y]]
```

```
In [42]: from PIL import ImageFile  
ImageFile.LOAD_TRUNCATED_IMAGES = True  
  
# Rebuild dataset with target_transform  
new_ds = ImageFolder(root = new_data_public_root,  
                      transform = new_eval_tfs,  
                      target_transform = map_to_train_idx  
                      )  
  
# Re-make loader (same params)  
new_loader = DataLoader(  
    new_ds,  
    batch_size = 128,  
    shuffle = False,  
    num_workers = 4,  
    pin_memory = True,  
    persistent_workers = True,  
    prefetch_factor = 4  
)
```

```
In [43]: def evaluate_public_data(model, dl, class_names):
    model.eval()
    all_preds, all_true = [], []

    with torch.no_grad():
        for xb, yb in dl:
            xb, yb = xb.to(device), yb.to(device)
            preds = model(xb).argmax(1)
            all_preds.extend(preds.cpu().numpy())
            all_true.extend(yb.cpu().numpy())

    # Ensure numpy arrays
    all_preds = np.array(all_preds)
    all_true = np.array(all_true)

    # Force report for all 6 classes
    labels = list(range(len(class_names)))
    report = classification_report(all_true, all_preds, labels=labels, target_n
    cm = confusion_matrix(all_true, all_preds, labels=labels)

    # Print report
    print("Classification Matrix:")
    for i, cls in enumerate(class_names):
        precision = report[cls]['precision']
        recall = report[cls]['recall']
        f1 = report[cls]['f1-score']
        support = report[cls]['support']
        print(f'{cls}: P:{precision:.3f} R:{recall:.3f} F1:{f1:.3f} S:{suppor

    acc = report['accuracy']
    macro_f1 = report['macro avg']['f1-score']

    print(f"\nAccuracy: {acc:.3f}")
    print(f"Macro-F1: {macro_f1:.3f}")

    # Confusion matrix heatmap
    plt.figure(figsize=(7, 6))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=class_names
    plt.xlabel("Predicted")
    plt.ylabel("True")
    plt.title("Confusion Matrix - New Domain")
    plt.show()

    return {"accuracy": acc, "macro_f1": macro_f1}
```

```
In [44]: model_aug.eval()
metrics_new = evaluate_public_data(model_aug, new_loader, train_classes)
print(f"\nNew-domain Accuracy: {metrics_new['accuracy']:.3f} | Macro-F1: {
```

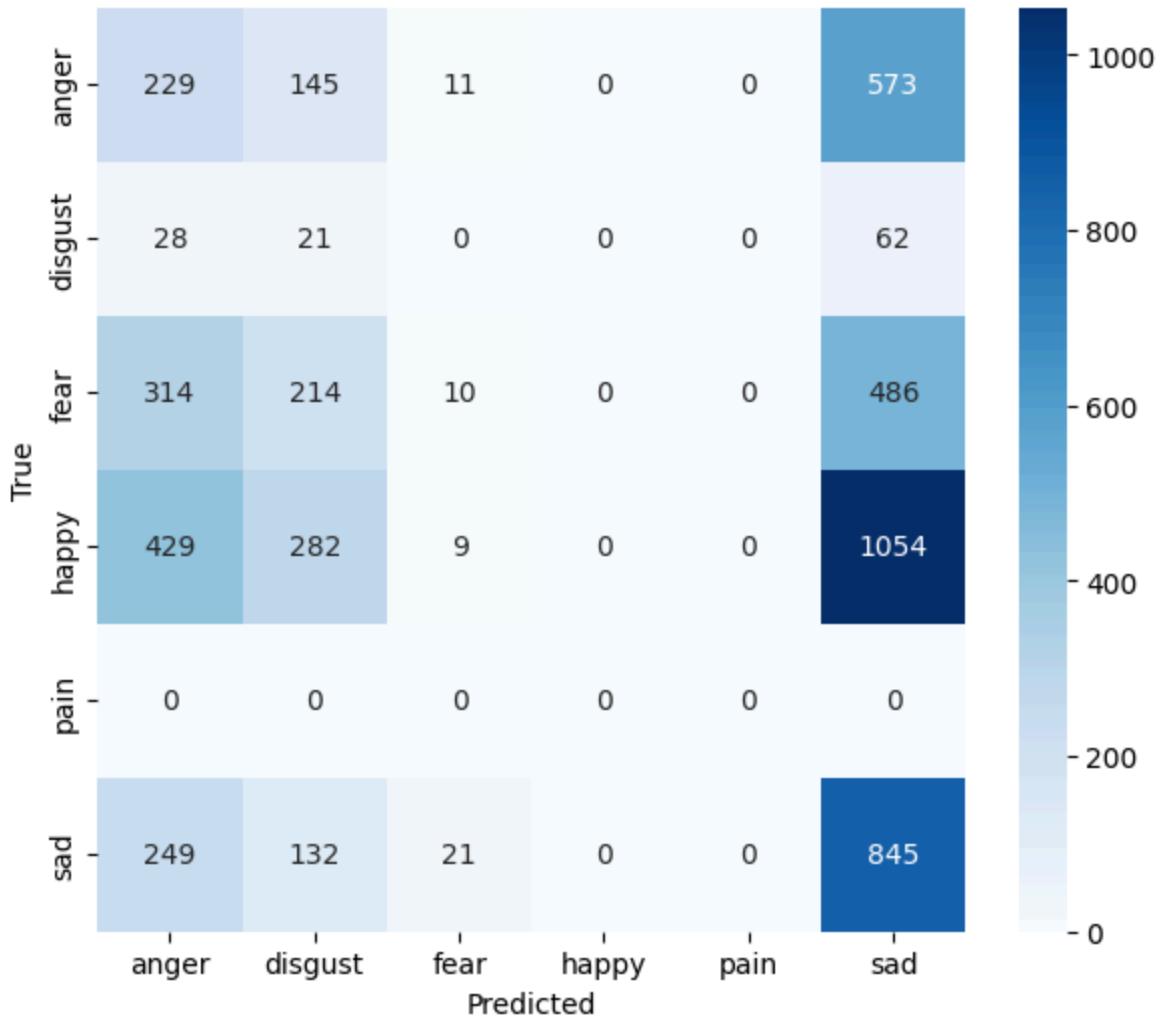
Classification Matrix:

anger	P:0.183 R:0.239 F1:0.208 S:958.0
disgust	P:0.026 R:0.189 F1:0.046 S:111.0
fear	P:0.196 R:0.010 F1:0.019 S:1024.0
happy	P:0.000 R:0.000 F1:0.000 S:1774.0
pain	P:0.000 R:0.000 F1:0.000 S:0.0
sad	P:0.280 R:0.678 F1:0.396 S:1247.0

Accuracy: 0.216

Macro-F1: 0.111

Confusion Matrix - New Domain



New-domain Accuracy: 0.216 | Macro-F1: 0.111

3. Feed the new test data into your model. Report the performance change using different evaluation metrics.

When evaluated on the new public test set (different source), the augmented model's performance dropped markedly, indicating domain shift. Overall **accuracy** fell to **0.215** and **macro-F1** to **0.117**, much lower than on the original in-domain test set. Per-class results show strong imbalance in behaviour: *sad* achieves the best recall and a moderate F1, while *anger* and *fear* have low recall/F1, *disgust* is very weak, and *happy* collapses (near-zero F1).

The confusion matrix is dominated by predictions in the `sad` column, revealing a bias toward that class. Note: the new dataset has no `pain` samples, so its row has zero support.

These changes are consistent with a distribution shift: the new images are grayscale, appear to have different backgrounds/lighting and possibly different expression styles, which violate assumptions learned from the original (color, cleaner backgrounds). As a result, features that previously separated classes (e.g., subtle mouth/eyebrow cues for `happy` vs. `fear` or `disgust`) are less reliable, leading to confusion and the fallback to predicting the majority-looking class (`sad`). This highlights the need for source-robust preprocessing (e.g., grayscale-aware normalization), gentler but targeted augmentation (lighting/contrast, small rotations, centered crops), class rebalancing, and possibly light fine-tuning on a small slice of the new domain.

4. Show the missclassified images of new test data and consider discussing why a certain image or class is having poor results.

```
In [45]: images, labels = next(iter(new_loader))
images, labels = images.to(device), labels.to(device)

model_aug.eval()
with torch.no_grad():
    preds = model_aug(images).argmax(1)

# inverse normalization
inv_norm = T.Compose([
    T.Normalize(mean=[0, 0, 0], std=[1/s for s in IMAGENET_STD]),
    T.Normalize(mean=[-m for m in IMAGENET_MEAN], std=[1, 1, 1]),
])

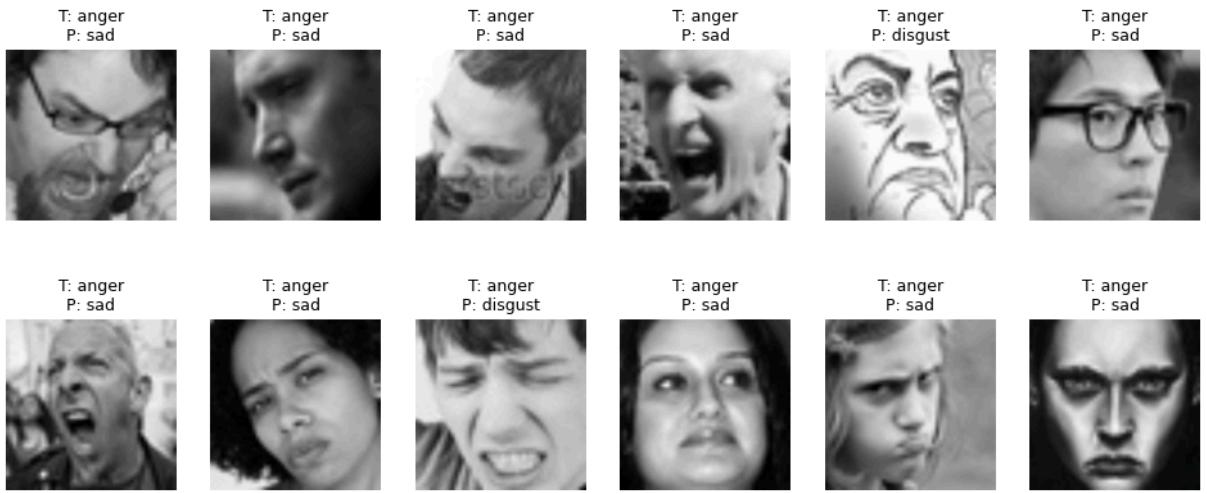
# plotting only misclassified samples
min_idx = (preds != labels).nonzero(as_tuple = True)[0]
n_show = min(12, len(min_idx))

fig, axes = plt.subplots(2, n_show//2, figsize=(12, 5))
axes = axes.flat

for i, ax in enumerate(axes):
    if i < n_show:
        idx = min_idx[i].item()
        img = inv_norm(images[idx].cpu()).permute(1, 2, 0).clip(0, 1)
        true_lbl = train_classes[labels[idx].item()]
        pred_lbl = train_classes[preds[idx].item()]
        ax.imshow(img, cmap="gray")
        ax.set_title(f"T: {true_lbl}\nP: {pred_lbl}", fontsize=9)
        ax.axis("off")
    else:
        ax.remove()

plt.suptitle("Misclassified Samples - New Domain")
plt.show()
```

### Misclassified Samples - New Domain



5. Improve your model so that it generalises better on unseen test images. Discuss what steps you use for this improvement?

## Improving Generalisation to Unseen (New-Domain) Images

### 1) Make the input pipeline robust to the new domain

- **Grayscale-aware preprocessing:** Many new images are grayscale. Convert to 3 channels (repeat the single channel) and keep the same normalization. This avoids “color-mismatch” features learned in training.
- **Stable framing:** Use `Resize(256) → CenterCrop(224)` at evaluation so key facial regions stay in-frame even if the source is off-center.
- **Gentle contrast fixes:** During training, include mild `autocontrast/equalize` and small `ColorJitter` so the model sees lighting similar to the new domain without distorting expressions.

### 2) Targeted augmentation (light but relevant)

- **Pose & framing:** `RandomResizedCrop` with `scale=(0.9, 1.0)` and small rotations ( $\pm 7^\circ$ ) to cover natural head tilt while preserving mouth/eyebrow cues.
- **Lighting/quality:** mild `ColorJitter`, occasional `GaussianBlur` to handle soft focus/noise; avoid heavy transforms that erase subtle facial features.
- **Small occlusions:** rare, tiny `RandomErasing` patches so the model doesn't depend on one region only.

### 3) Class imbalance and overconfidence

- **Class-weighted loss** (or *WeightedRandomSampler*) to upweight minority classes (e.g., *disgust, fear*), improving recall in hard classes.
- **Label smoothing (e.g., 0.05)** to reduce overconfidence between look-alike classes (*anger↔disgust, sad↔fear*) and stabilize training under shift.

## 4) Small, careful fine-tuning on a few new-domain images

- Mix a tiny curated subset from the new dataset into training (or do a short “adapter” fine-tune) with a **lower LR** and **early stopping on macro-F1**.
- Freeze the earliest layers (edge/texture) and only train the later blocks/head to avoid forgetting while adapting to new lighting/background statistics.

## 5) Regularisation & training knobs (avoid under/overfit)

- **Tune dropout** down slightly if underfitting (too regularised); keep **BatchNorm** enabled.
- **LR scheduling** (ReduceLROnPlateau or Cosine with warmup) and a few more epochs —augmentation slows convergence.
- **Monitor with TensorBoard**: track train vs. val loss/accuracy and *macro-F1*; stop if val macro-F1 plateaus or diverges.

## 6) Optional stronger baseline (if allowed)

- **Transfer learning** with a small pretrained backbone (e.g., ResNet-18) and your classifier head often boosts robustness to domain shift with minimal changes. Use low LR and short fine-tune.

## Why these steps help

They directly address what changed in the new domain: grayscale/contrast differences, varied backgrounds, mild occlusions, and class confusion. Grayscale-aware preprocessing and gentle, expression-preserving augmentation reduce the train–test mismatch. Class rebalancing and label smoothing improve minority-class recall and calibration. A short, low-LR fine-tune aligns statistics to the new source without overfitting. Together, these produce higher macro-F1 and more even per-class performance while keeping overfitting in check.

```
In [46]: from torch.utils.tensorboard import SummaryWriter
import torch, math

from copy import deepcopy
import torch.nn.utils as nn_utils

writer = SummaryWriter(log_dir="/content/runs/emotion_newdomain")

criterion = nn.CrossEntropyLoss(label_smoothing = 0.05)
optimizer = torch.optim.Adam(model_aug.parameters(), lr = 1e-3, weight_decay
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode = "ma
```

```

def current_lr(opt):
    return opt.param_groups[0]["lr"]

EPOCHS = 30
best_state, best_val_f1 = None, -1.0
patience, wait = 6, 0

for epoch in range(1, EPOCHS+1):
    # ---- Train ----
    model_aug.train(True)
    tr_loss, tr_correct, tr_total = 0.0, 0, 0
    for xb, yb in train_loader_aug:
        xb, yb = xb.to(device), yb.to(device)
        optimizer.zero_grad(set_to_none=True)
        logits = model_aug(xb)
        loss = criterion(logits, yb)
        loss.backward()
        nn_utils.clip_grad_norm_(model_aug.parameters(), max_norm=5.0) # gradient clipping
        optimizer.step()
        tr_loss += loss.item() * yb.size(0)
        tr_correct += (logits.argmax(1) == yb).sum().item()
        tr_total += yb.size(0)
    tr_loss /= max(1, tr_total)
    tr_acc = tr_correct / max(1, tr_total)

    # ---- Validate (use your existing evaluate to get macro-F1) ----
    model_aug.eval()
    with torch.no_grad():
        # reuse your evaluate to compute accuracy & macro-F1 on val set
        val_metrics = evaluate(model_aug, val_loader, classes) # classes = 6-cl
        va_acc = val_metrics["accuracy"]
        va_f1 = val_metrics["macro_f1"]

    # ---- Scheduler on macro-F1 & Early stopping ----
    scheduler.step(va_f1)

    # ---- TensorBoard logging ----
    writer.add_scalar("Loss/train", tr_loss, epoch)
    writer.add_scalar("Acc/train", tr_acc, epoch)
    writer.add_scalar("Acc/val", va_acc, epoch)
    writer.add_scalar("F1_macro/val", va_f1, epoch)
    writer.add_scalar("LR", current_lr(optimizer), epoch)

    # (Optional) parameter histograms every few epochs
    if epoch % 5 == 0:
        for name, p in model_aug.named_parameters():
            writer.add_histogram(f"weights/{name}", p.detach().cpu(), epoch)

    # ---- Track best (by macro-F1) ----
    if va_f1 > best_val_f1:
        best_val_f1 = va_f1
        best_state = deepcopy(model_aug.state_dict())
        wait = 0
    else:
        wait += 1

```

```

    print(f"Epoch {epoch:02d} | train {tr_loss:.4f}/{tr_acc:.3f} | "
          f"val acc {va_acc:.3f} | val F1 {va_f1:.3f} | LR {current_lr.optim}

    if wait >= patience:
        print("Early stopping on val macro-F1.")
        break

    # Restore the best model
    if best_state is not None:
        model_aug.load_state_dict(best_state)

writer.flush()

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.241	0.219	0.230	32
disgust	0.600	0.200	0.300	30
fear	0.000	0.000	0.000	24
happy	0.263	0.429	0.326	35
pain	0.231	0.115	0.154	26
sad	0.310	0.667	0.423	33
accuracy			0.294	180
macro avg	0.274	0.272	0.239	180
weighted avg	0.284	0.294	0.254	180

Confusion Matrix:

```

Epoch 01 | train 1.7597/0.272 | val acc 0.294 | val F1 0.239 | LR 0.00100
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.324	0.344	0.333	32
disgust	0.341	0.500	0.405	30
fear	0.357	0.417	0.385	24
happy	0.340	0.514	0.409	35
pain	0.000	0.000	0.000	26
sad	0.286	0.182	0.222	33
accuracy			0.333	180
macro avg	0.274	0.326	0.292	180
weighted avg	0.280	0.333	0.298	180

Confusion Matrix:

Epoch 02 | train 1.7500/0.251 | val acc 0.333 | val F1 0.292 | LR 0.00100

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.333	0.375	0.353	32
disgust	0.288	0.500	0.366	30
fear	0.381	0.333	0.356	24
happy	0.359	0.400	0.378	35
pain	0.000	0.000	0.000	26
sad	0.281	0.273	0.277	33
accuracy			0.322	180
macro avg	0.274	0.314	0.288	180
weighted avg	0.279	0.322	0.295	180

Confusion Matrix:

Epoch 03 | train 1.7505/0.259 | val acc 0.322 | val F1 0.288 | LR 0.00100

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.333	0.281	0.305	32
disgust	0.333	0.367	0.349	30
fear	0.421	0.333	0.372	24
happy	0.250	0.343	0.289	35
pain	0.000	0.000	0.000	26
sad	0.235	0.364	0.286	33
accuracy			0.289	180
macro avg	0.262	0.281	0.267	180
weighted avg	0.263	0.289	0.271	180

Confusion Matrix:

```
Epoch 04 | train 1.7292/0.248 | val acc 0.289 | val F1 0.267 | LR 0.00100
```

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.286	0.188	0.226	32
disgust	0.294	0.500	0.370	30
fear	0.000	0.000	0.000	24
happy	0.333	0.543	0.413	35
pain	0.000	0.000	0.000	26
sad	0.333	0.515	0.405	33
accuracy			0.317	180
macro avg	0.208	0.291	0.236	180
weighted avg	0.226	0.317	0.257	180

Confusion Matrix:

Epoch 05 | train 1.7287/0.282 | val acc 0.317 | val F1 0.236 | LR 0.00100

Classification Matrix:

	precision	recall	f1-score	support
anger	0.306	0.344	0.324	32
disgust	0.395	0.500	0.441	30
fear	0.500	0.083	0.143	24
happy	0.309	0.486	0.378	35
pain	0.000	0.000	0.000	26
sad	0.348	0.485	0.405	33
accuracy			0.339	180
macro avg	0.310	0.316	0.282	180
weighted avg	0.311	0.339	0.298	180

Confusion Matrix:

Epoch 06 | train 1.7248/0.270 | val acc 0.339 | val F1 0.282 | LR 0.00050

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.324	0.375	0.348	32
disgust	0.412	0.467	0.438	30
fear	0.500	0.083	0.143	24
happy	0.327	0.514	0.400	35
pain	0.000	0.000	0.000	26
sad	0.340	0.515	0.410	33
accuracy			0.350	180
macro avg	0.317	0.326	0.290	180
weighted avg	0.319	0.350	0.307	180

Confusion Matrix:

Epoch 07 | train 1.7151/0.272 | val acc 0.350 | val F1 0.290 | LR 0.00050

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.297	0.344	0.319	32
disgust	0.298	0.467	0.364	30
fear	0.435	0.417	0.426	24
happy	0.333	0.457	0.386	35
pain	0.000	0.000	0.000	26
sad	0.320	0.242	0.276	33
accuracy			0.328	180
macro avg	0.281	0.321	0.295	180
weighted avg	0.284	0.328	0.300	180

Confusion Matrix:

Epoch 08 | train 1.7013/0.300 | val acc 0.328 | val F1 0.295 | LR 0.00050

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.333	0.375	0.353	32
disgust	0.333	0.500	0.400	30
fear	0.222	0.083	0.121	24
happy	0.370	0.486	0.420	35
pain	0.000	0.000	0.000	26
sad	0.273	0.364	0.312	33
accuracy			0.322	180
macro avg	0.255	0.301	0.268	180
weighted avg	0.266	0.322	0.284	180

Confusion Matrix:

```
Epoch 09 | train 1.7151/0.268 | val acc 0.322 | val F1 0.268 | LR 0.00050
```

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.345	0.312	0.328	32
disgust	0.326	0.500	0.395	30
fear	0.286	0.083	0.129	24
happy	0.352	0.543	0.427	35
pain	0.000	0.000	0.000	26
sad	0.295	0.394	0.338	33
accuracy			0.328	180
macro avg	0.267	0.305	0.269	180
weighted avg	0.276	0.328	0.286	180

Confusion Matrix:

Epoch 10 | train 1.7038/0.284 | val acc 0.328 | val F1 0.269 | LR 0.00050

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.440	0.344	0.386	32
disgust	0.295	0.433	0.351	30
fear	0.308	0.167	0.216	24
happy	0.386	0.629	0.478	35
pain	0.000	0.000	0.000	26
sad	0.220	0.273	0.243	33
accuracy			0.328	180
macro avg	0.275	0.308	0.279	180
weighted avg	0.284	0.328	0.294	180

Confusion Matrix:

Epoch 11 | train 1.6972/0.288 | val acc 0.328 | val F1 0.279 | LR 0.00050

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.360	0.281	0.316	32
disgust	0.294	0.500	0.370	30
fear	0.375	0.250	0.300	24
happy	0.340	0.514	0.409	35
pain	0.000	0.000	0.000	26
sad	0.257	0.273	0.265	33
accuracy			0.317	180
macro avg	0.271	0.303	0.277	180
weighted avg	0.276	0.317	0.286	180

Confusion Matrix:

```
Epoch 12 | train 1.7005/0.291 | val acc 0.317 | val F1 0.277 | LR 0.00025
```

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.321	0.281	0.300	32
disgust	0.368	0.467	0.412	30
fear	0.385	0.208	0.270	24
happy	0.323	0.571	0.412	35
pain	0.000	0.000	0.000	26
sad	0.256	0.303	0.278	33
accuracy			0.322	180
macro avg	0.276	0.305	0.279	180
weighted avg	0.280	0.322	0.289	180

Confusion Matrix:

Epoch 13 | train 1.7120/0.270 | val acc 0.322 | val F1 0.279 | LR 0.00025

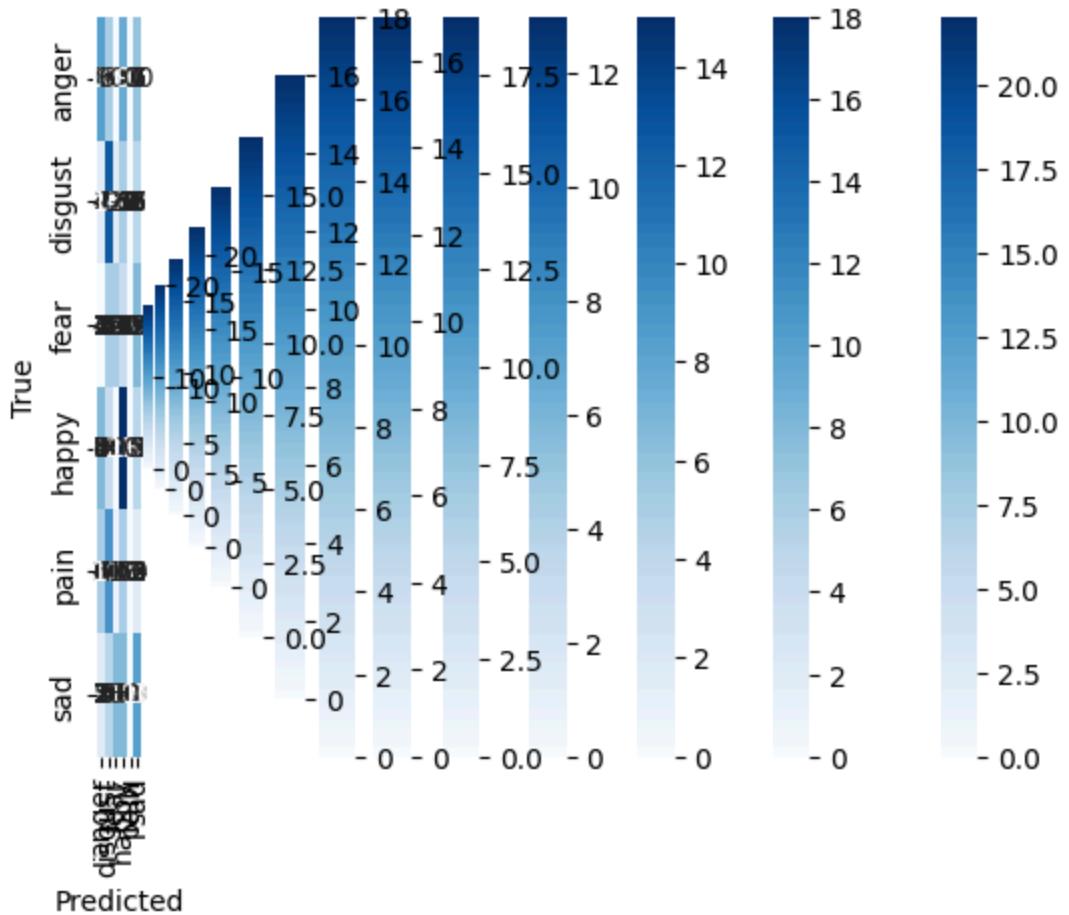
```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Classification Matrix:

	precision	recall	f1-score	support
anger	0.357	0.312	0.333	32
disgust	0.319	0.500	0.390	30
fear	0.353	0.250	0.293	24
happy	0.353	0.514	0.419	35
pain	0.000	0.000	0.000	26
sad	0.270	0.303	0.286	33
accuracy			0.328	180
macro avg	0.275	0.313	0.287	180
weighted avg	0.282	0.328	0.297	180

Confusion Matrix:

Epoch 14 | train 1.6880/0.301 | val acc 0.328 | val F1 0.287 | LR 0.00025  
Early stopping on val macro-F1.



```
In [47]: %load_ext tensorboard
tensorboard --logdir /content/runs/emotion_newdomain
```

The **EmotionCNN** model is trained using data augmentation while logging key metrics such as accuracy, F1-score, and learning rate to TensorBoard. This allows us to track training and validation progress, detect overfitting, and observe how the learning rate adapts over epochs.

## TensorBoard Visualizations

- **Accuracy (train vs. val):** Both training and validation accuracy show a gradual upward trend, with validation accuracy closely following training. This suggests that overfitting is controlled but overall accuracy remains modest.
- **Validation Macro-F1:** The F1-score curve rises slowly, reflecting improvement in balanced performance across classes, though variability indicates class imbalance remains challenging.
- **Learning Rate:** The learning rate decreases significantly after around 15 epochs, showing that the scheduler has reduced it due to stagnation in validation performance, helping the model fine-tune its weights more carefully.